

Q1

Chosen Parameters


Below are the chosen μ and Σ values chosen for the 4 gaussian class conditionals

$$\mu_1 = [1 \ 0 \ 0] \quad \mu_2 = [0 \ 0 \ 2.5] \quad \mu_3 = [0 \ 0.5 \ 0] \quad \mu_4 = [1 \ 0 \ 10]$$

$$\Sigma_1 = I^{3 \times 3} \quad \Sigma_2 = I^{3 \times 3} \quad \Sigma_3 = I^{3 \times 3} \quad \Sigma_4 = I^{3 \times 3}$$


The below data distribution results in an optimal MAP classifier accuracy of ~80% using true knowledge of the data when the classifier is applied to validation set with 100K samples.

Data distribution




Kfold cross validation for perceptron selection

From the below figure, kfold generally favours higher perceptron counts at lower dataset sizes. At high dataset sizes it favours lower perceptron counts. Unfortunately multiple kfold runs weren't performed multiple times for each dataset, nor were enough datasets generated with kfold run on them to establish a clear trend. A violation of the above assumption is the optimally chosen 11 perceptrons for the 500 dataset. Observing the histograms of accuracy vs perceptron count for each dataset can shed more light on this unexpected behaviour. In the following figures, it will become apparent that the optimally selected perceptron counts rely on minute differences based on average trained model performance on validation partitions during kfold. These differences are so minute they can almost be disregarded which renders the below graph less useful than expected.




Dataset 100

Perceptron accuracy doesn't really plateau at any particular perceptron value. There is a clear trend though of increased accuracy vs perceptron count. Some perceptrons experience counts result in convergence issues where the minimizer gets stuck at a local minimum at perceptron. This can be seen at perceptron count = 16. Low perceptron counts result in generally bad model performance.




Dataset 500

Perceptron accuracy plateaus at around 8 perceptrons with negligible increase in accuracy at higher perceptron counts. Because all models beyond 8 perceptrons here are very close to the aspirational performance level established by the optimal MAP classifier, the model selected by kfold is only marginally better than the model at the accuracy plateau in the below figure. Additionally, the accuracy from which kfold selects the optimal model is calculated from the mean of the classification accuracy of the models when applied to 1 fold of the training dataset. For small data sizes, that accuracy measurement may not be representative enough of model performance to select the optimal perceptron count.




Dataset 1000

Perceptron accuracy plateaus at around 5 perceptrons with negligible increase in accuracy at higher perceptron counts. We are beginning to see a trend here of lower perceptron counts being sufficient for near optimal accuracy performance. More data with fewer perceptrons results in models that can generalize more effectively than more complicated models trained with fewer data that might overfit. Additionally, the 1k dataset with kfold has a larger validation partition to work with when selecting the optimal perceptron count.




Dataset 5000

Perceptron accuracy plateaus at around 4 perceptrons with negligible increase in accuracy at higher perceptron counts. This observation can be used to establish a trend of MLPs befitting from larger datasets while requiring smaller models to achieve adequate performance. Higher perceptron counts beyond 20 have not been explored but a downwards trend in terms of accuracy is to be expected due to likely overfitting of the model.



Accuracy vs Dataset Size

From the below graph it is clear that more data for MLPs (regardless of model complexity) results in better performing models. Increasing model complexity must be justified by greater data availability. Otherwise overly complex models are unnecessarily computationally expensive and can overfit to the available data.



Code Discussion

A few notes concerning the implementation of Q1:

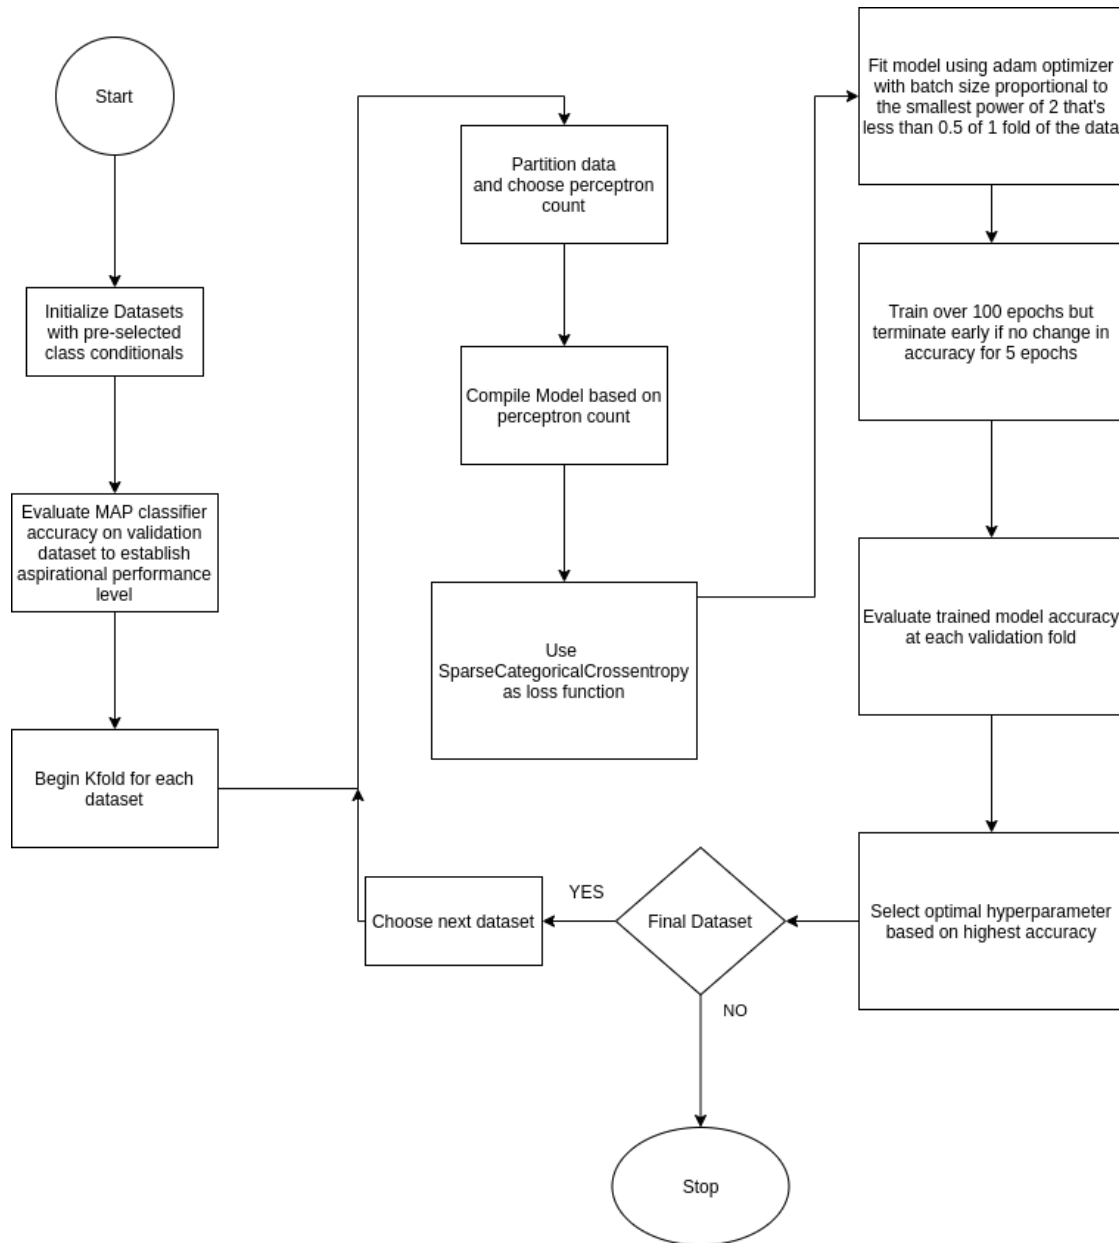
- TensorFlow with Cuda acceleration was used for this question
- SparseCategoricalLoss was used instead of categorical loss because the label data was not encoded using 1 hot encoding and was instead encoded with integers from 0 to 3
- The models output layer was normalized with softmax so the loss functions from_logits flag was set to False because the outputs of the network are not logits
- Softplus activation was used for the hidden layer as an alternative to Smooth-ReLU
- Training max epochs was set to 100 with early termination using keras callbacks depending on changes in accuracy. If no change in accuracy (subject to a default delta) is observed for 5 epochs the training terminates early and returns the last trained model
- Batch sizes for training were selected based on the following equation:

$2^{\log(\text{floor}((\text{dataset_size}/\text{kfolds})*0.50))}$ i.e. select closes power of 2 that's smaller than $\frac{1}{2}$ the data

This was done to keep batch sizes to multiples of 2 to speed up communication between CPU and GPU.

Code FlowChart

Below is a general flow chart of the code for this question. Please note that the code was split across two files. Q1.py and analysis.py. Q1 produces the required data for this question and analysis generates the required graphs available above in this report.



Q2

Chosen Parameters

The parameters: \mathbf{a} and $\mathbf{\Sigma}$ are randomized everytime the code is run. The below results were generated from the following parameters

$$\mathbf{a} = [56.5067, 65.1923, 23.8135, 72.2253, 41.4401, 62.4820, 77.4701]$$

$$\mathbf{\Sigma} = 0.5114 \mathbf{I}$$

Noise variance and L2 Regularizer fluctuation as follows:

$$\in [10^{-3}, 10^{-3}]$$

$$\in [10^{-3}, 10^{-2}]$$

MAP parameter estimation derivation

$$y = \omega^T x + \omega_0 + v$$

$$y = \theta^T z + v \quad v \sim \mathcal{N}(0, 1)$$

$\omega \in \mathbb{R}^d$

prior $\theta \in \mathbb{R}^{d+1}$ $\theta = \begin{bmatrix} \omega_0 \\ \omega \end{bmatrix} \quad z = \begin{bmatrix} 1 \\ x \end{bmatrix}$

$$\theta \sim \mathcal{N}(0, \beta I)$$

$$\hat{\theta}_{\text{map}} = \underset{\theta}{\text{argmax}} \ln p(\theta | D)$$

where $D = \{(x_1, y_1), (x_2, y_2), \dots\}$
Assuming iid samples

$$\begin{aligned} \hat{\theta}_{\text{map}} &= \underset{\theta}{\text{argmax}} \ln p(D | \theta) + \ln p(\theta) \\ &= \underset{\theta}{\text{argmax}} \ln p(\theta) + \sum_{i=1}^N \ln p(x_i, y_i | \omega) \\ &= \underset{\theta}{\text{argmax}} \ln p(\theta) + \sum_{i=1}^N \ln (p(y_i | \omega, x_i) p(x_i | \omega)) \\ &\quad \text{Assume } x_i \perp \omega \quad \xrightarrow{\text{Drop this}} \\ &= \underset{\theta}{\text{argmax}} \ln p(\theta) + \sum_{i=1}^N \ln (p(y_i | \omega, x_i)) \end{aligned}$$

$$y \sim N(\theta^T z, 1)$$

$$\ln p(\theta) = \ln \left((2\pi)^{\frac{cd+1}{2}} |\beta \Sigma|^{-\frac{1}{2}} e^{-\frac{1}{2}(\omega-\theta)^T (\beta \Sigma)^{-1}(\omega-\theta)} \right)$$

$$= \underbrace{\ln \left((2\pi)^{-\frac{d+1}{2}} |\beta I|^{-\frac{1}{2}} \right)}_{\text{Constant w.r.t } \theta} - \frac{1}{2\beta} \omega^T \omega$$

$$\ln(p(y|x, \omega)) = \ln \left((2\pi)^{-\frac{1}{2}} e^{-\frac{(y-\theta^T z)^2}{2}} \right)$$

$$= \underbrace{-\frac{1}{2} \ln(2\pi) - \frac{1}{2} (y - \theta^T z)^2}_{\text{Constant w.r.t } \theta}$$

$$\hat{\theta}_{MAP} = \underset{\theta}{\operatorname{argmax}} \frac{1}{2\beta} \theta^T \theta + \sum_{i=1}^N \frac{1}{2} (y_i - \theta^T z_i)^2$$

$$= \underset{\theta}{\operatorname{argmin}} \frac{1}{\beta} \theta^T \theta + \sum_{i=1}^N (y_i - \theta^T z_i)^2 \quad \text{[multiplied by } 2 \text{]}$$

$$\frac{\partial}{\partial \theta^T} \left(\sum_{i=1}^N (y_i - \theta^T z_i)^2 + \frac{1}{\beta} \theta^T \theta \right)$$

$$= \sum_{i=1}^N \frac{\partial}{\partial \theta^T} (y_i - \theta^T z_i)^2 + \frac{2\theta}{\beta}$$

$$= - \sum_{i=1}^N 2(y_i - \theta^T z_i) z_i + \frac{2\theta}{\beta}$$

$$= - \sum y_i z_i + \sum \theta^T z_i z_i + \frac{1}{\beta} \theta$$

$$= - \sum y_i z_i + \sum z_i z_i^T \hat{\theta}_{MAP} + \frac{1}{\beta} \hat{\theta}_{MAP} = 0$$

$$\sum z_i z_i^T \theta + \frac{1}{\beta} \theta = \sum y_i z_i$$

$$\left(\frac{1}{N} \sum z_i z_i^T + \frac{1}{\beta N} I \right) \hat{\theta}_{MAP} = \frac{1}{N} \sum y_i z_i$$

$$R(\beta)$$

$$q$$

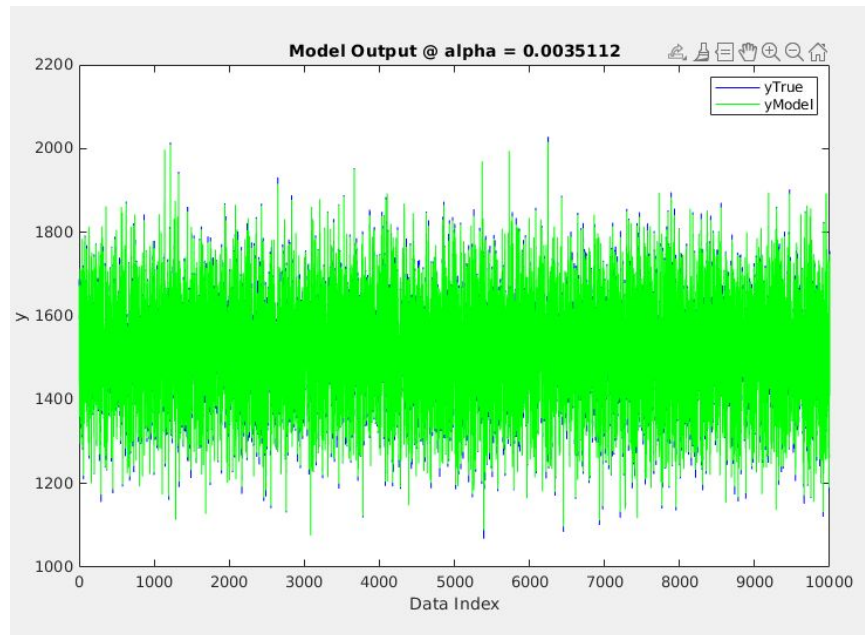
$$\xrightarrow{\quad} \hat{\theta}_{MAP} = R(\beta)^{-1} q \xleftarrow{\quad}$$

$\lim_{\beta \rightarrow \infty} \hat{\theta}_{MAP} = \hat{\theta}_{ML}$ because θ prior becomes irrelevant.

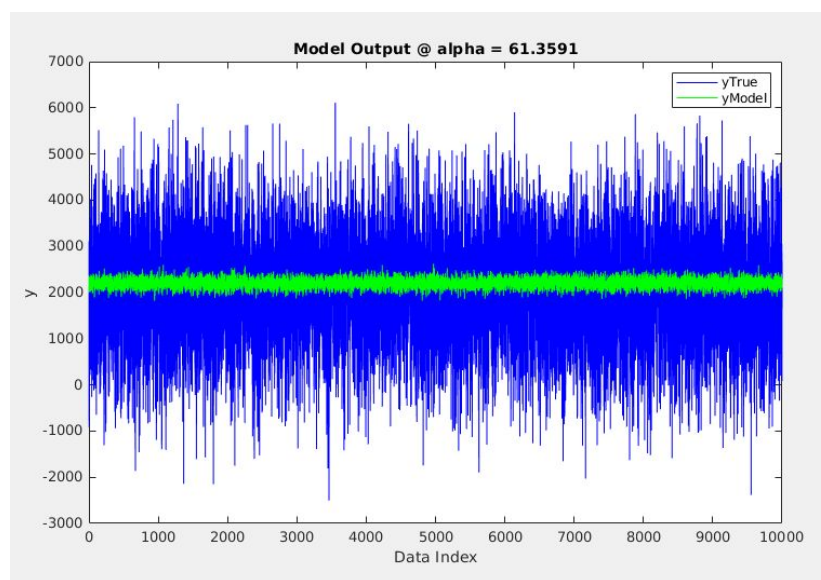
Results

Verifying derivation

At low alpha levels the derived MAP estimator manages to predict the data with high accuracy as shown in the below graph. The achieved MSE for this particular run was 93.4474.




If noise is allowed to dominate the model the the MAP estimator fails to find an appropriate fit. Introducing excess noise affects optimal L2 regularizer selection which will be discussed in the next section.



Results

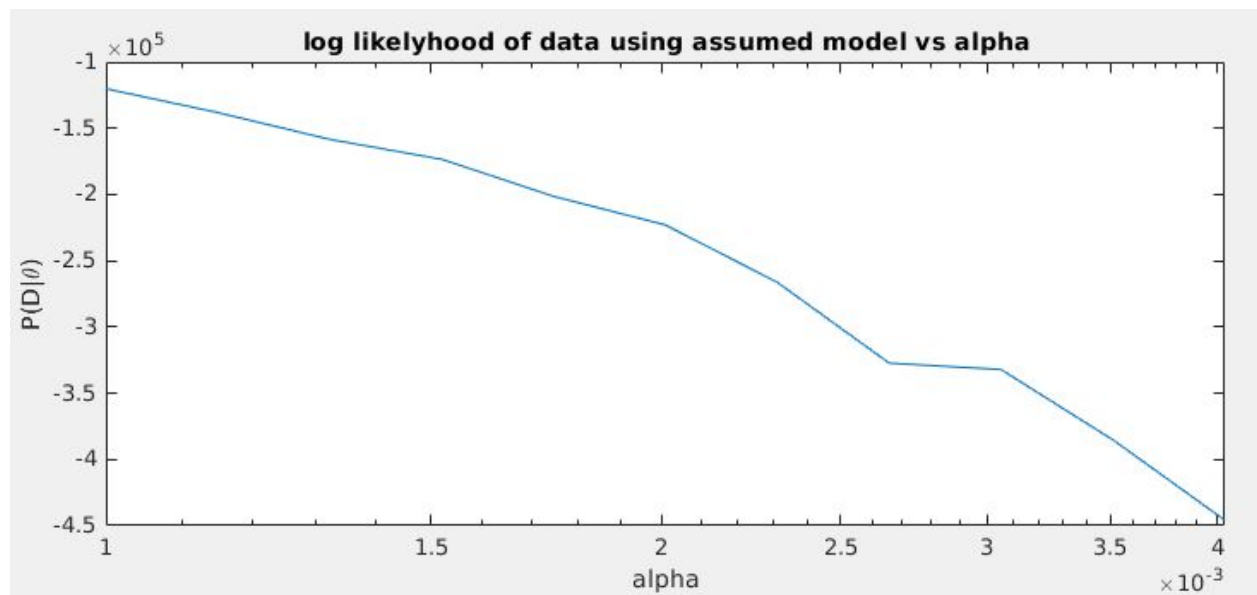
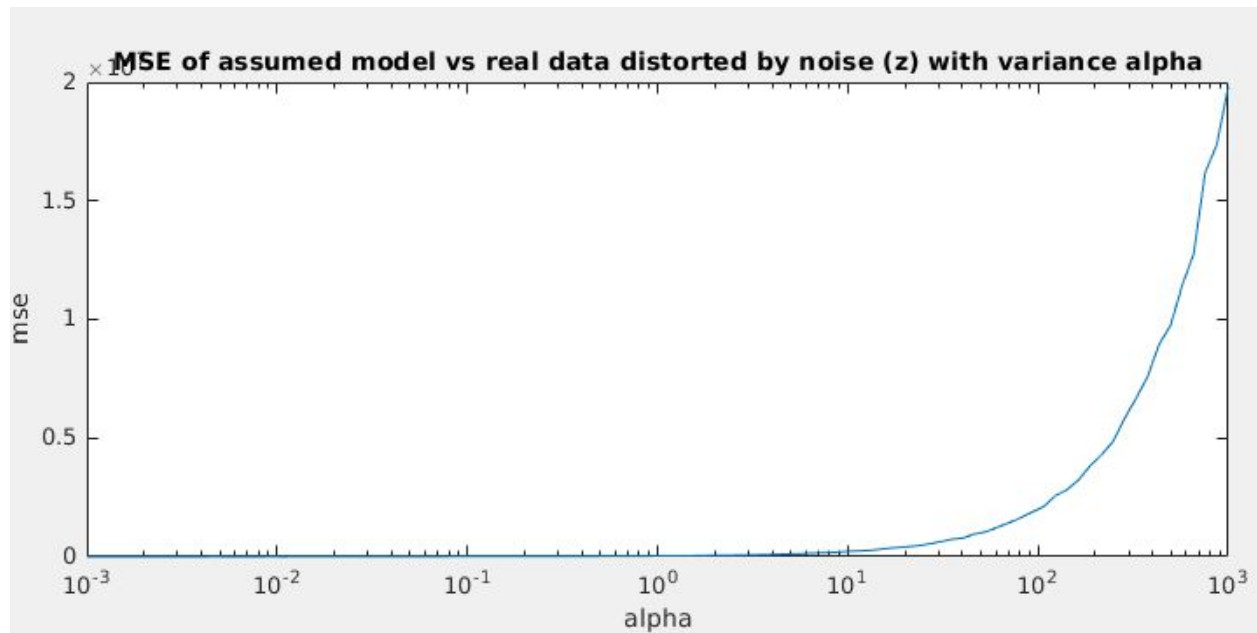
Noise and Regularization

At lower noise levels applied to x , the optimally chosen β tends towards higher values thus MAP parameter estimation approaches ML parameter estimation which results in the prior on the weights becoming irrelevant. This is understandable because, at low noise levels, the assumed model is close to the actual model of the data. Increasing noise, k fold selects smaller β s to combat the increased variance of the data at high noise levels until eventually it always selects the lowest β available. At high enough noise levels, noise completely dominates the output of the model rendering regression meaningless.



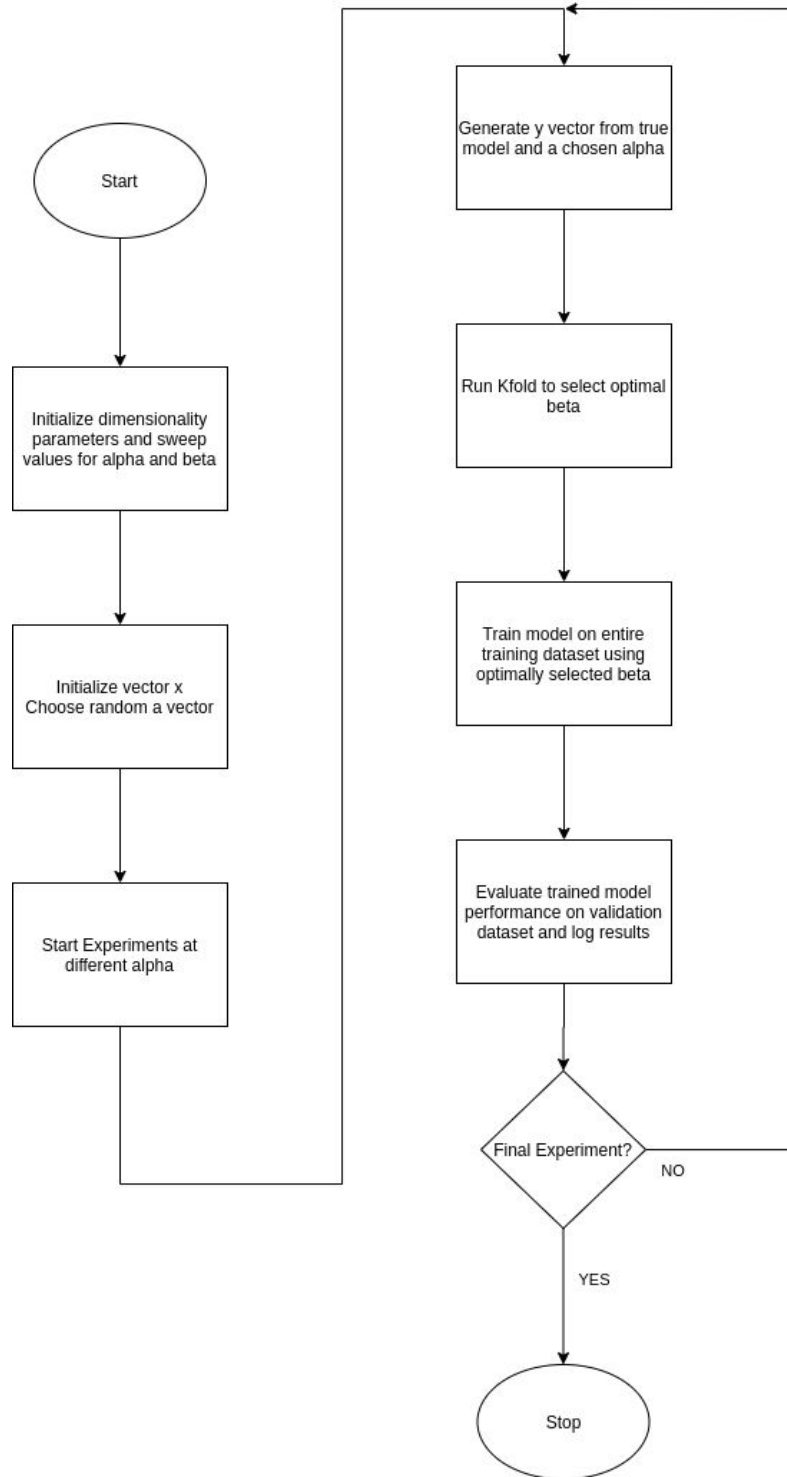
Data likelihood vs noise

After the optimal beta is selected during kfold, MSE and data likelihood are evaluated on the validation dataset, the below figures show the result of the evaluation at each noise level applied to the data. As expected MSE and data likelihood are heavily affected by the noise introduced in the data. MSE increases sharply at high enough noise variance and data likelihood falls accordingly regardless of the chosen beta.



Code discussion

The overall flow of the code is given in the below figure



Code also available at:

https://github.com/asultan123/ML_HW3

Q1 MLP Training Code

```
import os
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
import numpy as np
from scipy.stats import multivariate_normal
import math
import random
from functools import partial
from multiprocessing import Pool
import pickle

class GuassainClassConditionalPdf:
    def __init__(self, mean, sigma):
        self.mean = mean
        self.sigma = sigma
    def generate(self, count):
        return np.random.multivariate_normal(self.mean, self.sigma, count)
    def evaluate(self, data):
        return multivariate_normal(mean=self.mean, cov=self.sigma).pdf(data);

class Data:
    def __init__(self, features = None, labels = None, count = 0, classCount = 0,
gaussianConditionals=None):
        self.features = features
        self.labels = labels
        self.count = count
        self.classCount = classCount
        self.gaussianConditionals = gaussianConditionals
    def generateData(self, gaussianConditionals, count, classPriors = None):
        self.count = count
        self.gaussianConditionals = gaussianConditionals
        self.classCount = len(gaussianConditionals)
```

```

    if classPriors is None:
        self.labels = np.array([math.floor(num) for num in
np.random.uniform(0,self.classCount,count)])
        self.labels.sort()
        self.labelCount = [np.count_nonzero(self.labels == l, axis=0) for l in range(0,
self.classCount)]
        self.features = []
        for eachClass in range(0, self.classCount):

self.features.extend(gaussianConditionals[eachClass].generate(self.labelCount[eachClass]).toli
st())

        self.features = np.array(self.features)
        data_aggregate = list(zip(self.labels, self.features))
        random.shuffle(data_aggregate)
        self.labels = np.array([label for label, feature in data_aggregate])
        self.features = np.array([np.array(feature) for label, feature in data_aggregate])
        self.labels = self.labels.astype('float32')
        self.features = self.features.astype('float32')

def plotData3D(self):
    fig = plt.figure()
    ax = fig.add_subplot(111, projection='3d')
    for l in range(0, self.classCount):
        labelIndexes = np.where(self.labels==l)[0]
        x1 = self.features[labelIndexes,0]
        x2 = self.features[labelIndexes,1]
        x3 = self.features[labelIndexes,2]
        ax.scatter(x1,x2,x3)
    plt.show()

def mapClassifyWithTrueData(self):
    pxGivenL = []
    for eachClass in range(0, self.classCount):
        pxGivenL.append(self.gaussianConditionals[eachClass].evaluate(self.features))
    classifyResult = np.array(np.argmax(np.array(pxGivenL), axis=0))
    labels = np.array(self.labels)
    accuracy = np.count_nonzero(np.equal(classifyResult, labels))
    return accuracy/len(labels)

def partition(self, partitionCount):
    elementsPerParition = math.floor(self.count/10)
    partitionedDataFeatures = np.array([self.features[partition:partition+elementsPerParition] for
partition in range(0, self.count, elementsPerParition)])
    partitionedDataLabels = np.array([self.labels[partition:partition+elementsPerParition] for
partition in range(0, self.count, elementsPerParition)])
    dataList = []

```

```

        for eachPartitionLabels, eachPartitionFeatures in zip(partitionedDataLabels,
partitionedDataFeatures):
            dataList.append(Data(eachPartitionFeatures, eachPartitionLabels, elementsPerPartition,
self.classCount, self.gaussianConditionals))
        return dataList
    @classmethod
    def fuse(cls, partitions):
        features = []
        labels = []
        for eachPartition in partitions:
            features.extend(eachPartition.features)
            labels.extend(eachPartition.labels)
        gaussians = partitions[0].gaussianConditionals
        classCount = len(gaussians)
        return Data(np.array(features), np.array(labels), len(features), classCount, gaussians)

```

```

class ModelHelper:
    @classmethod
    def createModel(cls, output_size, perceptrons):
        model = tf.keras.models.Sequential([
            tf.keras.layers.Dense(perceptrons, activation='softplus'),
            tf.keras.layers.Dense(output_size, activation='softmax'),
        ])
        loss_fn = tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False)
        model.compile(optimizer='adam', loss=loss_fn, metrics=['accuracy'])
        return model
    @classmethod
    def fitModel(cls, batch_size, model, data, monitor='accuracy', min_delta=0.001, patience=5,
epochs=100):
        callback = tf.keras.callbacks.EarlyStopping(monitor=monitor, min_delta=min_delta,
patience=patience)
        model.fit(np.array(data.features, dtype=np.float32), np.array(data.labels, dtype=np.float32),
epochs=epochs, batch_size=batch_size, verbose = 0 , callbacks=[callback])
        return model
    @classmethod
    def evalAccuracy(cls, model, data):
        predictions = model.predict(data.features).numpy()
        classifyResult = np.array(np.argmax(predictions, axis=1))
        return (np.count_nonzero(classifyResult==data.labels))/data.count

def kfoldCrossValidation(kfolds, dataset, modelCreator, modelFitter, modelEvaluator,
hyperparamater_range):
    def keywithmaxval(d):
        v=list(d.values())

```

```

k=list(d.keys())
return k[v.index(max(v))]

hyperparameter_accuracy = {}
for hyperparameter in hyperparameter_range:
    print('Hyperparameter = ' + str(hyperparameter))
    model = modelCreator(hyperparameter)
    partitionedDataset = dataset.partition(kfolds)
    accuracy = [0]*kfolds
    for k in range(0, kfolds):
        print('K = ' + str(k))
        trainingData = Data.fuse(partitionedDataset[0:k] +
partitionedDataset[k+1:len(partitionedDataset)])
        validationData = partitionedDataset[k:k+1][0]
        fittedModel = modelFitter(model, trainingData)
        accuracy[k] = modelEvaluator(model, validationData)
    hyperparameter_accuracy[hyperparameter] = np.mean(accuracy)

optimal_hyperparameter = keywithmaxval(hyperparameter_accuracy)
optimal_model = modelCreator(optimal_hyperparameter)
fitted_optimal_model = modelFitter(optimal_model, dataset)

return (optimal_hyperparameter, hyperparameter_accuracy, fitted_optimal_model)

if __name__ == "__main__":
    print(tf.test.is_gpu_available())
    d1 = Data()
    d2 = Data()
    d3 = Data()
    d4 = Data()
    dvalid = Data()
    c1 = GuassianClassConditionalPdf([1,0,0], [[1,0,0],[0,1,0],[0,0,1]])
    c2 = GuassianClassConditionalPdf([0,0,2.5], [[1,0,0],[0,1,0],[0,0,1]])
    c3 = GuassianClassConditionalPdf([0,0.5,0], [[1,0,0],[0,1,0],[0,0,1]])
    c4 = GuassianClassConditionalPdf([1,0,10], [[1,0,0],[0,1,0],[0,0,1]])
    d1.generateData([c1, c2, c3, c4], 100)
    d2.generateData([c1, c2, c3, c4], 500)
    d3.generateData([c1, c2, c3, c4], 1000)
    d4.generateData([c1, c2, c3, c4], 5000)
    datasets = [d1, d2, d3, d4]
    dvalid.generateData([c1, c2, c3, c4], 100000)

kfolds = 10

```

```

results = {}
for eachDataset in datasets:
    print('Dataset: ' + str(eachDataset.count))

    results[str(eachDataset.count)] = {}
    results[str(eachDataset.count)]['map_accuracy'] =
eachDataset.mapClassifyWithTrueData()
    results[str(eachDataset.count)]['dataset'] = {}
    results[str(eachDataset.count)]['optimal_hyperparameter'] = {}
    results[str(eachDataset.count)]['hyperParameter_accuracy'] = {}
    results[str(eachDataset.count)]['fittedModel_validation_accuracy'] = {}

    modelCreator = partial(ModelHelper.createModel, 4)
    batch_size = 2**int(math.log((eachDataset.count/kfolds)*0.50,2))
    print('Dataset batch size: ' + str(batch_size))

    modelFitter = partial(ModelHelper.fitModel, batch_size)

    (optimal_hyperparameter, hyperParameter_accuracy, fitted_optimal_model) =
kfoldCrossValidation(kfolds, eachDataset, modelCreator, modelFitter,
ModelHelper.evalAccuracy, hyperparameter_range = range(1,5))
    fittedModel_validation_accuracy = ModelHelper.evalAccuracy(fitted_optimal_model, dvalid)

    results[str(eachDataset.count)]['dataset'] = eachDataset
    results[str(eachDataset.count)]['optimal_hyperparameter'] = optimal_hyperparameter
    results[str(eachDataset.count)]['hyperParameter_accuracy'] = hyperParameter_accuracy
    results[str(eachDataset.count)]['fittedModel_validation_accuracy'] =
fittedModel_validation_accuracy
    with open('TEST_final_results.pkl', 'wb') as output:
        pickle.dump(results, output, pickle.HIGHEST_PROTOCOL)

print('Done')

```

Q1 Analysis Code

```
import matplotlib.pyplot as plt
import numpy as np
from scipy.stats import multivariate_normal
import math
import random
from functools import partial
from multiprocessing import Pool
import pickle
import pprint
```

```
class GuassianClassConditionalPdf:
    def __init__(self, mean, sigma):
        self.mean = mean
        self.sigma = sigma
    def generate(self, count):
        return np.random.multivariate_normal(self.mean, self.sigma, count)
    def evaluate(self, data):
        return multivariate_normal(mean=self.mean, cov=self.sigma).pdf(data);
```

```
class Data:
    def __init__(self, features = None, labels = None, count = 0, classCount = 0,
gaussianConditionals=None):
        self.features = features
        self.labels = labels
        self.count = count
        self.classCount = classCount
        self.gaussianConditionals = gaussianConditionals
    def generateData(self, gaussianConditionals, count, classPriors = None):
        self.count = count
        self.gaussianConditionals = gaussianConditionals
        self.classCount = len(gaussianConditionals)
        if classPriors is None:
            self.labels = np.array([math.floor(num) for num in
np.random.uniform(0,self.classCount,count)])
```

```

        self.labels.sort()
        self.labelCount = [np.count_nonzero(self.labels == l, axis=0) for l in range(0,
self.classCount)]
        self.features = []
        for eachClass in range(0, self.classCount):

self.features.extend(gaussianConditionals[eachClass].generate(self.labelCount[eachClass]).toli
st())
        self.features = np.array(self.features)
        data_aggregate = list(zip(self.labels, self.features))
        random.shuffle(data_aggregate)
        self.labels = np.array([label for label, feature in data_aggregate])
        self.features = np.array([np.array(feature) for label, feature in data_aggregate])
        self.labels = self.labels.astype('float32')
        self.features = self.features.astype('float32')

def plotData3D(self):
    fig = plt.figure()
    ax = fig.add_subplot(111, projection='3d')
    for l in range(0, self.classCount):
        labelIndexes = np.where(self.labels==l)[0]
        x1 = self.features[labelIndexes,0]
        x2 = self.features[labelIndexes,1]
        x3 = self.features[labelIndexes,2]
        ax.scatter(x1,x2,x3)

    ax.set_xlabel('X1 Label')
    ax.set_ylabel('X2 Label')
    ax.set_zlabel('X3 Label')
    ax.set_title('Data distribution')
    plt.show()
    return fig

def mapClassifyWithTrueData(self):
    pxGivenL = []
    for eachClass in range(0, self.classCount):
        pxGivenL.append(self.gaussianConditionals[eachClass].evaluate(self.features))
    classifyResult = np.array(np.argmax(np.array(pxGivenL), axis=0))
    labels = np.array(self.labels)
    accuracy = np.count_nonzero(np.equal(classifyResult, labels))
    return accuracy/len(labels)

def partition(self, partitionCount):
    elementsPerPartition = math.floor(self.count/10)

```



```

        partitionedDataFeatures = np.array([self.features[partition:partition+elementsPerPartition] for
partition in range(0, self.count, elementsPerPartition)])
        partitionedDataLabels = np.array([self.labels[partition:partition+elementsPerPartition] for
partition in range(0, self.count, elementsPerPartition)])
        dataList = []
        for eachPartitionLabels, eachPartitionFeatures in zip(partitionedDataLabels,
partitionedDataFeatures):
            dataList.append(Data(eachPartitionFeatures, eachPartitionLabels, elementsPerPartition,
self.classCount, self.gaussianConditionals))
        return dataList
    @classmethod
    def fuse(cls, partitions):
        features = []
        labels = []
        for eachPartition in partitions:
            features.extend(eachPartition.features)
            labels.extend(eachPartition.labels)
        gaussians = partitions[0].gaussianConditionals
        classCount = len(gaussians)
        return Data(np.array(features), np.array(labels), len(features), classCount, gaussians)

```

class ModelHelper:

```

    @classmethod
    def createModel(cls, output_size, perceptrons):
        model = tf.keras.models.Sequential([
            tf.keras.layers.Dense(perceptrons, activation='softplus'),
            tf.keras.layers.Dense(output_size, activation='softmax'),
        ])
        loss_fn = tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False)
        model.compile(optimizer='adam', loss=loss_fn, metrics=['accuracy'])
        return model
    @classmethod
    def fitModel(cls, batch_size, model, data, monitor='accuracy', min_delta=0.001, patience=5,
epochs=100):
        callback = tf.keras.callbacks.EarlyStopping(monitor=monitor, min_delta=min_delta,
patience=patience)
        model.fit(np.array(data.features, dtype=np.float32), np.array(data.labels, dtype=np.float32),
epochs=epochs, batch_size=batch_size, verbose = 0 , callbacks=[callback])
        return model
    @classmethod
    def evalAccuracy(cls, model, data):
        predictions = model.predict(data.features).numpy()
        classifyResult = np.array(np.argmax(predictions, axis=1))
        return (np.count_nonzero(classifyResult==data.labels))/data.count

```

```

if __name__ == "__main__":
    with open('final_results.pkl', 'rb') as input:
        results = pickle.load(input)
        pp = pprint.PrettyPrinter(indent=4)
        pp.pprint(results)

    # plot 5k dataset
    fig = results['5000']['dataset'].plotData3D()

    # plot histograms of perceptron accuracies
    for dataset_key, dataset_results in results.items():
        fig = plt.figure()
        plt.bar(range(1,20), list(dataset_results['hyperParameter_accuracy'].values()), color='blue')
        plt.xlabel('Perceptrons')
        plt.ylabel('Model accuracy')
        plt.title('MLP chosen perceptrons @ data set size = ' + dataset_key)
        plt.show()

    # Plot optimal perceptrons
    chosenPerceptrons = []
    datasetSizes = []
    for dataset_key, dataset_results in results.items():
        chosenPerceptrons.append(dataset_results['optimal_hyperparameter'])
        datasetSizes.append(dataset_results['dataset'].count)
    print(datasetSizes)
    print(chosenPerceptrons)
    plt.bar(datasetSizes, chosenPerceptrons, width=100, color='b')
    plt.xlabel('Dataset Sizes')
    plt.ylabel('Optimal Perceptron Count')
    plt.xticks(datasetSizes)
    plt.yticks(chosenPerceptrons)
    plt.show()

    # Map Accuracy vs Optimal NN accuracy
    map_accuracy = []
    optimal_NN_accuracy = []
    datasetSizes = []
    gaussians = results['5000']['dataset'].gaussianConditionals
    dValid = Data()
    dValid.generateData(gaussians, 100000)
    for dataset_key, dataset_results in results.items():
        map_accuracy.append(dValid.mapClassifyWithTrueData())
        datasetSizes.append(dataset_results['dataset'].count)

```

```

    optimal_NN_accuracy.append(dataset_results['fittedModel_validation_accuracy'])
plt.plot(datasetSizes,map_accuracy, color='g', label='MAP Classifier with knowledge')
plt.plot(datasetSizes,optimal_NN_accuracy, color='r', label='Optimally selected NN Classifier')
plt.xlabel('Dataset Size')
plt.ylabel('Optimal Model Accuracy')
plt.gca().set_yticklabels(['{:.2f}%'.format(x*100) for x in plt.gca().get_yticks()])
plt.legend()
plt.show()

```

Q2 Code

clear all, close all

% ===== Parameter Setup =====%

```

n=7;
Ntrain=100;
Ntest=10000;
lower_noise_multiplier = 10^-3;
upper_noise_multiplier = 10^3;
x_mu_max=10;
totalExp = 100;
betasPerExperiment = 100;
maxExpPerJob = 10;
maxJobs = floor(totalExp/maxExpPerJob);
betaMin = -3;
betaMax = 2;

```

% ===== Data Generation =====%

```

[x, gm] = generateRandGMMDData(1, n, Ntrain+Ntest, x_mu_max);
wtrue = rand(n,1)*100;

% alpha = trace(gm.Sigma(:,1))/n*10.^(linspace(-2,2, totalExp));
alpha = logspace(-3,3, totalExp);
expRange = floor(linspace(0,totalExp-1,maxExpPerJob));

chosenBetas = zeros(1,totalExp);
mse = zeros(1,totalExp);
data_likelihoood = zeros(1,totalExp);

```

```

% ===== Experiment Runs at Increasing Alphas =====%

for exp = 1:totalExp

exp,

% ===== Generate Noisy Data =====%

[y, x] = noisy_model(n, Ntrain+Ntest, wtrue, x, alpha(exp));
data = segment_data(y, x, Ntrain);

% ===== Use Kfold to optimize L2 Regularizer for Wmap =====%

[bestBeta, wfit] = runfoldexperiment_ridge_regression_bestbeta(Ntrain, 10, betaMin, betaMax,
betasPerExperiment, data.train, 1);

experiment{exp}.data = data;
chosenBetas(exp) = bestBeta;
chosenWfits{exp} = wfit;

% ===== Evaluate MSE for optimally L2 Regularizer =====%

[~, total_mse, ~] = calc_MSE_noisy_linear_model(data.test.y,data.test.x,wfit);
mse(exp) = total_mse;

data_likelyhood(exp) = data_log_likelyhood(data.test.y,data.test.x, wfit, 1);
end

figure(1)
subplot(3,1,1), semilogx(alpha, chosenBetas)
title('Chosen weight regularization value beta for noise (z) variance alpha');
xlabel('alpha'),ylabel('beta'),
hold on, subplot(3,1,2), semilogx(alpha, mse)
title('MSE of assumed model vs real data distorted by noise (z) with variance alpha');
xlabel('alpha'),ylabel('mse'),
hold on, subplot(3,1,3), semilogx(alpha, data_likelyhood)
title('log likelyhood of data using assumed model vs alpha');
xlabel('alpha'),ylabel('P(D|{\theta})'),

figure(2),
subplot(2,1,1),
plot_experiment_data(10, alpha, experiment, chosenWfits)
title(['Model Output @ alpha = ', num2str(alpha(10))]);

```

```
xlabel('Data Index'),ylabel('y'), legend('yTrue', 'yModel');
```

```
subplot(2,1,2),  
plot_experiment_data(80, alpha, experiment, chosenWfits)  
title(['Model Output @ alpha = ', num2str(alpha(80))]);  
xlabel('Data Index'),ylabel('y'), legend('yTrue', 'yModel');
```

```
function plot_experiment_data(exp, alpha, experiment, chosenWfits)  
noise_alpha = alpha(exp);  
noise_alpha  
[pred, total_mse, ~] =  
calc_MSE_noisy_linear_model(experiment{exp}.data.test.y,experiment{exp}.data.test.x,chosen  
Wfits{exp});  
total_mse  
plot_data_and_noise_model_pred(experiment{exp}.data.test.y, pred)  
end
```

```
function [pred, total_mse, mse] = calc_MSE_noisy_linear_model(y,x,wfit)  
N = size(x,2);  
b = [ones(1,N); x];  
pred = wfit*b;  
mse = (y-wfit*b).^2;  
total_mse = mean(mse);  
end
```

```
function [bestBeta, wfit] = runfoldexperiment_ridgeregression_bestbeta(N, K, betaMin,  
betaMax, betaCount, dataset, model_noise_variance)  
partitions_idx_start = ceil(linspace(0,N,K+1));  
indPartitionLimits = zeros(K, 2);  
for k = 1:K  
indPartitionLimits(k,:) = [partitions_idx_start(k)+1,partitions_idx_start(k+1)];  
end  
loglikelyood_Beta = zeros(1,betaCount);  
beta = 10.^linspace(betaMin, betaMax, betaCount);  
for betaldx = 1:betaCount  
loglikelihood = zeros(1,K);  
for k= 1:K  
indValidate = indPartitionLimits(k,1):indPartitionLimits(k,2);  
x = dataset.x;  
y = dataset.y;  
xValidate = x(:, indValidate); % Using folk k as validation set  
yValidate = y(:, indValidate);  
if k == 1
```

```

        indTrain = indPartitionLimits(k+1,1):N;
    elseif k == K
        indTrain = 1:indPartitionLimits(k-1,2);
    else
        indTrain = [1:indPartitionLimits(k-1,2),indPartitionLimits(k+1,1):N];
    end
    xTrain = x(:, indTrain); % using all other folds as training set
    yTrain = y(:, indTrain); % using all other folds as training set

    wfit = fit_noisy_linear_model(yTrain, xTrain, beta(betaldx), model_noise_variance);
%     [~, total_mse, ~] = calc_MSE_noisy_linear_model(yValidate,xValidate,wfit);
%     loglikelihood(k) = total_mse;

    loglikelihood(k) = data_log_likelihood(yValidate,xValidate, wfit, model_noise_variance);
end
loglikelihood_Beta(betaldx) = mean(loglikelihood);
end
 [~, bestBetaldx] = max(loglikelihood_Beta);
%     [~, bestBetaldx] = min(loglikelihood_Beta);
    bestBeta = beta(bestBetaldx);
    wfit = fit_noisy_linear_model(dataset.y, dataset.x, bestBeta, model_noise_variance);
end

function plot_data_and_noise_model_pred(y, pred)
    plot(y,'b'), hold on
    plot(pred, 'g'),
end

function data = segment_data(y, x, Ntrain)
    data.train.x = x(:, 1:Ntrain);
    data.train.y = y(:, 1:Ntrain);
    data.test.x = x(:, Ntrain+1:end);
    data.test.y = y(:, Ntrain+1:end);
end

function [y, x] = noisy_model(n, N, wtrue, x, alpha)
    z = generateRandWhiteNoise(1, n, N, alpha);
    v = generateRandWhiteNoise(1, 1, N, 1);
    y = wtrue.*(x+z)+v;
end

function log_likelihood = data_log_likelihood(y, x, w, noise_variance)
    N = size(x,2);
    b = [ones(1,N); x];

```

```

% log_likelihood =
sum(log((1/(sqrt(2*pi)*noise_variance))*exp(-(y-w*b)/(2*(noise_variance)^2))));
likelyhoods = zeros(1,N);
mus = w*b;
for i = 1:N
    likelyhoods(i) = evalGaussian(y(i), mus(i), noise_variance.^2);
end
log_likelihood = sum(log(likelyhoods));
end

```

```

function g = evalGaussian(x,mu,Sigma)
% Evaluates the Gaussian pdf N(mu,Sigma) at each column of X
[n,N] = size(x);
C = ((2*pi)^n * det(Sigma))^(1/2);
E = -0.5*sum((x-repmat(mu,1,N)).*(inv(Sigma)*(x-repmat(mu,1,N))),1);
g = C*exp(E);
end

```

```

function wfit = fit_noisy_linear_model(y,x, beta, noise_variance)
N = size(x,2);
dim = size(x,1)+1;
b = [ones(1,N); x];
y = repmat(y, dim, 1);
R = b*b';
Q = sum(y.*b, 2);
wfit = inv(R+(noise_variance^2/beta).*eye(dim))*Q;
end

```

```

function [data, gmdist] = generateRandWhiteNoise(Order, dim, N, alpha)
sigma = zeros(dim,dim,Order);
mu = zeros(Order, dim);
for m = 1:Order
    sigma(:,:,m) = alpha.*eye(dim);
end
gmdist = gmdistribution(mu,sigma,ones(1,Order)/Order);
data = random(gmdist,N);
end

```

```

function [data, gmdist] = generateRandGMMData(Order, dim, N, mu_max)
sigma = zeros(dim,dim,Order);
mu = zeros(Order, dim);
for m = 1:Order
    sigma(:,:,m) = rand(1,1).*eye(dim);
    mu(m, :) = rand(1,dim)*mu_max;
end

```

```
end
gmdist = gmdistribution(mu,sigma,ones(1,Order)/Order);
data = random(gmdist,N);
end
```